



General Tree & Indexing

LI Hao 李颢, Assoc. Prof. SPEIT & Dept. Automation of SEIEE

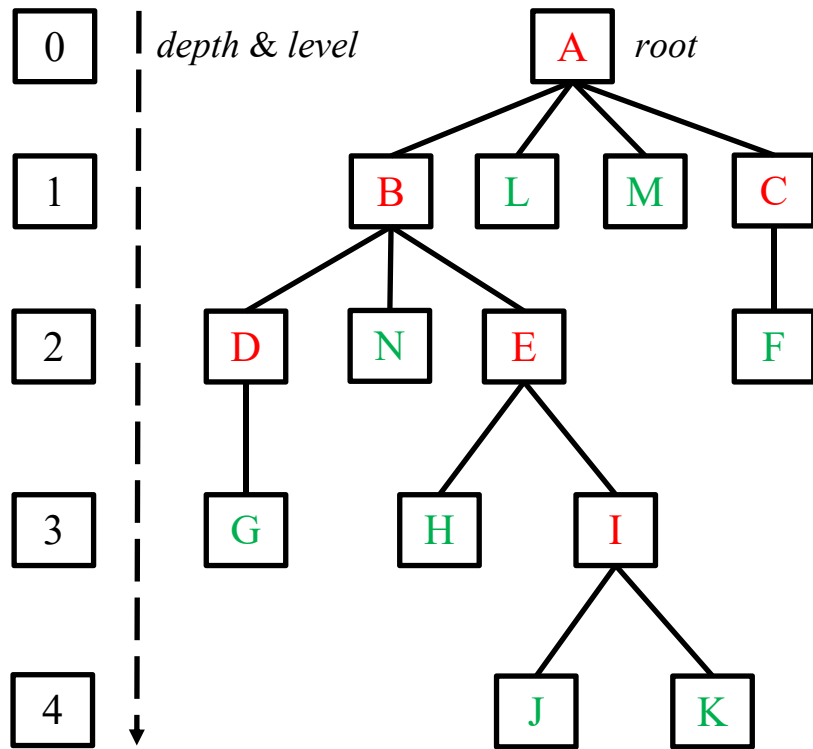


上海交通大学
SHANGHAI JIAO TONG UNIVERSITY

General Tree

- General tree

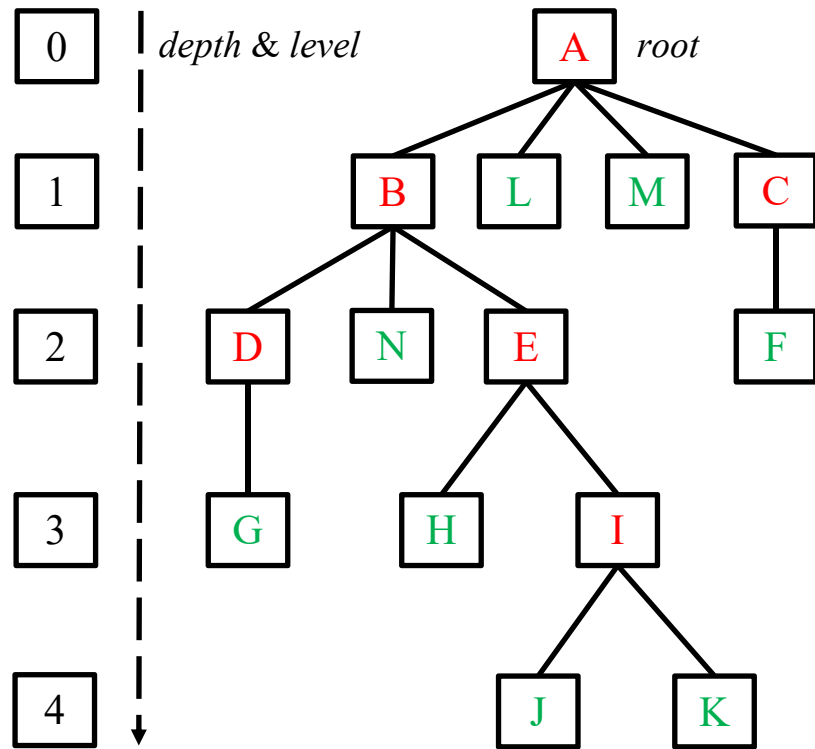
- node & edge
- root
- arbitrary number of subtrees
- parent & children (not limited to 2)
- ancestor & descendant
- path & length
 - B-E-I, A-C-F : two paths of length 2
 - A-B-E-I-J: a path of length 4
- depth (cardinal) & level (ordinal)
- height (largest depth+1)
- leaf node & internal node



General Tree

- General tree


- node & edge
- root
- arbitrary number of subtrees
- parent & children (not limited to 2)
- out degree - number of children
- left-most child - arranged from left to right
- ancestor & descendant
- path & length
- depth (cardinal) & level (ordinal)
- height (largest depth+1)
- leaf node & internal node



General Tree



- General tree
 - left-child/right-sibling* implementation




```

template <class T> class GTNode{ // GT (general tree) node abstract class
private:T e; // node's element
    GTNode *cL,*sR,*p; // node's leftmost child, right sibling, parent
public: GTNode(){cL=sR=p=NULL;} ~GTNode(){}
    GTNode(const T& ei,GTNode* cLi=NULL,GTNode* sRi=NULL,GTNode* pi=NULL){
        e=ei;cL=cLi;sR=sRi;p=pi;}
    const T& getE() const{return e;} void setE(const T& ei){e=ei;}
    inline GTNode* getC() const{return cL;} void setC(GTNode* g){cL=g;}
    inline GTNode* getS() const{return sR;} void setS(GTNode* g){sR=g;}
    inline GTNode* getP() const{return p;} void setP(GTNode* g){p=g;}
    bool isLeaf(){return cL==NULL;}
};

inline void indent(int L){while(L-->0) std::cout<<"|...";}
// preorder traversal of GT
template <class T> void showSubGT(GTNode<T>* g,int L){
    if(g==NULL) return;indent(L);std::cout<<g->getE()<<"\n";
    g=g->getC();while(g!=NULL){showSubGT(g,L+1);g=g->getS();}}
template <class T> void showGT(GTNode<T>* g){
    while(g->getP()!=NULL) g=g->getP(); showSubGT(g,0);}
    
```


General Tree

- General tree
 - left-child/right-sibling* implementation



sub-GT of B =>	GT to which B belongs =>
B	A
...D	...B
... ...GD
...NG
...EN
... ...HE
... ...IH
... JI
... KJ
K
	...L
	...M
	...C
F


```

GTNode<char> *A=new GTNode<char>('A');GTNode<char> *B=new GTNode<char>('B');GTNode<char> *C=new GTNode<char>('C');
GTNode<char> *D=new GTNode<char>('D');GTNode<char> *E=new GTNode<char>('E');GTNode<char> *F=new GTNode<char>('F');
GTNode<char> *G=new GTNode<char>('G');GTNode<char> *H=new GTNode<char>('H');GTNode<char> *I=new GTNode<char>('I');
GTNode<char> *J=new GTNode<char>('J');GTNode<char> *K=new GTNode<char>('K');GTNode<char> *L=new GTNode<char>('L');
GTNode<char> *M=new GTNode<char>('M');GTNode<char> *N=new GTNode<char>('N');
A->setC(B);B->setP(A);B->setS(L);B->setC(D);C->setP(A);C->setC(F);D->setP(B);D->setS(N);D->setC(G);
E->setP(B);E->setC(H);F->setP(C);G->setP(D);H->setP(E);H->setS(I);I->setP(E);I->setC(J);J->setP(I);
J->setS(K);K->setP(I);L->setP(A);L->setS(M);M->setP(A);M->setS(C);N->setP(B);N->setS(E);
cout<<"sub-GT of B => \n";showSubGT<char>(B,0);cout<<"GT to which B belongs => \n";showGT<char>(B);
    
```

General Tree



- General tree
 - left-child/right-sibling* implementation




```

template <class T> class LGT{ // linked general tree abstract class
private:GTNode<T>* r; // GT root
    inline void indent(int L){while(L-->0) std::cout<<"|---";}
    void SIn(GTNode<T>* g,int L);
    void clearIn(GTNode<T>* g);
public: LGT(const T& ei){r=new GTNode<T>(ei);} ~LGT(){
    inline GTNode<T>* getR() const{return r;}
    void setC(LGT* gt){r->setC(gt->getR());delete gt;} // merge 'gt' as left-most child & deallocate it
    void setS(LGT* gt){r->setS(gt->getR());delete gt;} // merge 'gt' as right sibling & deallocate it
    // setC() & setS() deallocate gt's LGT space but not gt's associated GTNodes space
    void clear(); // deallocate root's associated GTNodes space
    void S(); // preorder traversal of GT
};

template <class T> void LGT<T>::SIn(GTNode<T>* g,int L){if(g==NULL) return;
    while(g!=NULL){indent(L);std::cout<<g->getE()<<"\n";SIn(g->getC(),L+1);g=g->getS();}}
// default arguments are only allowed in function declarations (so omitted in S() definition here)
template <class T> void LGT<T>::S(){SIn(r,0);}
template <class T> void LGT<T>::clearIn(GTNode<T>* g){if(g==NULL) return;
    GTNode<T>* t=g->getC();delete g;g=t;while(g!=NULL){t=g->getS();clearIn(g);g=t;}}
template <class T> void LGT<T>::clear(){clearIn(r);r=NULL;}
    
```

General Tree

- General tree
 - left-child/right-sibling* implementation



```


GT of A =>
A
|
|---B
|---|---D
|---|---|---G
|---|---N
|---|---E
|---|---|---H
|---|---|---I
|---|---|---J
|---|---|---K
|---L
|---M
|---C
|---|---F
GT of A after clearing =>
    
```

```

LGT<char> *gtA=new LGT<char>('A');LGT<char> *gtB=new LGT<char>('B');LGT<char> *gtC=new LGT<char>('C');
LGT<char> *gtD=new LGT<char>('D');LGT<char> *gtE=new LGT<char>('E');LGT<char> *gtF=new LGT<char>('F');
LGT<char> *gtG=new LGT<char>('G');LGT<char> *gtH=new LGT<char>('H');LGT<char> *gtI=new LGT<char>('I');
LGT<char> *gtJ=new LGT<char>('J');LGT<char> *gtK=new LGT<char>('K');LGT<char> *gtL=new LGT<char>('L');
LGT<char> *gtM=new LGT<char>('M');LGT<char> *gtN=new LGT<char>('N');
gtJ->setS(gtK);gtI->setC(gtJ);gtH->setS(gtI);gtE->setC(gtH);gtN->setS(gtE);gtD->setC(gtG);gtD->setS(gtN);
gtB->setC(gtD);gtC->setC(gtF);gtM->setS(gtC);gtL->setS(gtM);gtB->setS(gtL);gtA->setC(gtB);
cout<<"GT of A =>\n";gtA->S();gtA->clear();cout<<"GT of A after clearing =>\n";gtA->S();
    
```


General Tree

- **General tree**
 - *parent pointer* implementation
 - determine if two nodes are in the same tree
 - e.g. $G \rightarrow D \rightarrow B$, $K \rightarrow I \rightarrow E \rightarrow B$, so G & K in same tree
 - e.g. $H \rightarrow E \rightarrow B$, $L \rightarrow A$, so H & L not in same tree
 - merge two trees




General Tree

- **General tree**
 - *parent pointer* implementation
 - determine if two nodes are in the same tree
 - e.g. $G \rightarrow D \rightarrow B$, $K \rightarrow I \rightarrow E \rightarrow B$, so G & K in same tree
 - e.g. $H \rightarrow E \rightarrow B$, $L \rightarrow A$, so H & L not in same tree
 - merge two trees
 - set a tree's root as parent of the other's root



General Tree

- General tree
 - parent pointer implementation




```
template <class T> class PPNode{ // parent pointer GT node abstract class
private:T e;int n; // node's element; number of nodes of the node's tree
    PPNode *p; // node's parent
public: PPNode(){p=NULL;n=1;} ~PPNode(){} PPNode(const T& ei){e=ei;p=NULL;n=1;}
    const T& getE() const{return e;} void setE(const T& ei){e=ei;}
    int getN() const{return n;} void setN(int ni){n=ni;}
    inline PPNode* getP() const{return p;} void setP(PPNode* g){p=g;}
    static PPNode* find(PPNode* g){while(g->getP()!=NULL) g=g->getP();return g;}
    static void ppunion(PPNode* a,PPNode* b){a=find(a);b=find(b);if(a==b) return;
        if(a->getN()<=b->getN()){a->setP(b);b->setN(b->getN()+a->getN());}
        else{b->setP(a);a->setN(a->getN()+b->getN());}}
};

// ostream overloading, so that 'cout<<{PPNode<T> object}' can have meaning
template <typename T> std::ostream& operator<<(std::ostream& out,PPNode<T>* b){
    out<<b->getE()<<':'<<b->getN(); return out;}

PPNode<char>* p1,*p2;PPNode<char>* pp[26]={new PPNode<char>('A'),
new PPNode<char>('B'), new PPNode<char>('C'), new PPNode<char>('D'), new PPNode<char>('E'), new PPNode<char>('F'),
new PPNode<char>('G'), new PPNode<char>('H'), new PPNode<char>('I'), new PPNode<char>('J'), new PPNode<char>('K'),
new PPNode<char>('L'), new PPNode<char>('M'), new PPNode<char>('N'), new PPNode<char>('O'), new PPNode<char>('P'),
new PPNode<char>('Q'), new PPNode<char>('R'), new PPNode<char>('S'), new PPNode<char>('T'), new PPNode<char>('U'),
new PPNode<char>('V'), new PPNode<char>('W'), new PPNode<char>('X'), new PPNode<char>('Y'), new PPNode<char>('Z')};
while(true){int i=rand()%26,j=rand()%26;p1=PPNode<char>::find(pp[i]);p2=PPNode<char>::find(pp[j]);
    cout<<"merge "<<pp[i]<<" & "<<pp[j]<<" : "<<pp[i]<<" => "<<p1<<" | "<<pp[j]<<" => "<<p2<<" : ";
    if(p1==p2){cout<<p1<<"=="<<p2<<" unnecessary to merge!\n";}
    else{PPNode<char>::ppunion(p1,p2);
        if(p1->getN()<p2->getN()){cout<<" merge into "<<p2<<endl;if(26==p2->getN()) break;}
        else{cout<<" merge into "<<p1<<endl;if(26==p1->getN()) break;}}}
```

General Tree

- General tree
 - parent pointer implementation



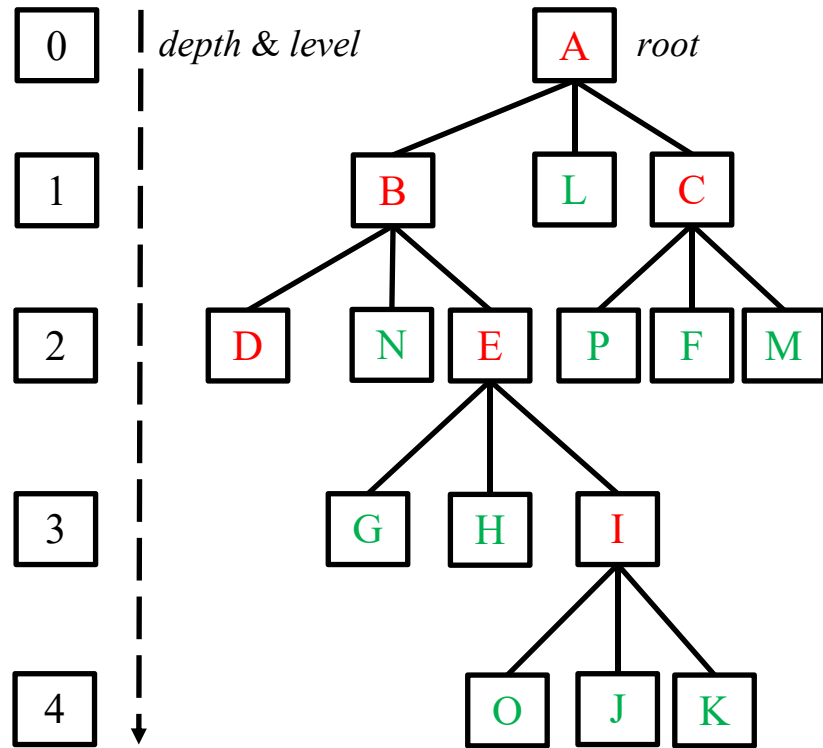
```

merge N:1 & W:1 : N:1 => N:1 | W:1 => W:1 : merge into W:2
merge L:1 & R:1 : L:1 => L:1 | R:1 => R:1 : merge into R:2
merge B:1 & B:1 : B:1 => B:1 | B:1 => B:1 : B:1==B:1 unnecessary to merge!
merge M:1 & Q:1 : M:1 => M:1 | Q:1 => Q:1 : merge into Q:2
merge B:1 & H:1 : B:1 => B:1 | H:1 => H:1 : merge into H:2
merge C:1 & D:1 : C:1 => C:1 | D:1 => D:1 : merge into D:2
merge A:1 & R:2 : A:1 => A:1 | R:2 => R:2 : merge into R:3
merge Z:1 & O:1 : Z:1 => Z:1 | O:1 => O:1 : merge into O:2
merge W:2 & K:1 : W:2 => W:2 | K:1 => K:1 : merge into W:3
merge K:1 & Y:1 : K:1 => W:3 | Y:1 => Y:1 : merge into W:4
merge H:2 & I:1 : H:2 => H:2 | I:1 => I:1 : merge into H:3
merge D:2 & D:2 : D:2 => D:2 | D:2 => D:2 : D:2==D:2 unnecessary to merge!
merge Q:2 & S:1 : Q:2 => Q:2 | S:1 => S:1 : merge into Q:3
merge C:1 & D:2 : C:1 => D:2 | D:2 => D:2 : D:2==D:2 unnecessary to merge!
merge X:1 & R:3 : X:1 => X:1 | R:3 => R:3 : merge into R:4
merge J:1 & M:1 : J:1 => J:1 | M:1 => Q:3 : merge into Q:4
merge O:2 & W:4 : O:2 => O:2 | W:4 => W:4 : merge into W:6
merge F:1 & R:4 : F:1 => F:1 | R:4 => R:4 : merge into R:5
merge X:1 & S:1 : X:1 => R:5 | S:1 => Q:4 : merge into R:9
merge J:1 & Y:1 : J:1 => R:9 | Y:1 => W:6 : merge into R:15
merge B:1 & L:1 : B:1 => H:3 | L:1 => R:15 : merge into R:18
merge D:2 & B:1 : D:2 => D:2 | B:1 => R:18 : merge into R:20
merge E:1 & F:1 : E:1 => E:1 | F:1 => R:20 : merge into R:21
merge S:1 & A:1 : S:1 => R:21 | A:1 => R:21 : R:21==R:21 unnecessary to merge!
merge R:21 & C:1 : R:21 => R:21 | C:1 => R:21 : R:21==R:21 unnecessary to merge!
merge B:1 & Y:1 : B:1 => R:21 | Y:1 => R:21 : R:21==R:21 unnecessary to merge!
merge N:1 & E:1 : N:1 => R:21 | E:1 => R:21 : R:21==R:21 unnecessary to merge!
merge C:1 & D:2 : C:1 => R:21 | D:2 => R:21 : R:21==R:21 unnecessary to merge!
merge Y:1 & G:1 : Y:1 => R:21 | G:1 => G:1 : merge into R:22
merge G:1 & X:1 : G:1 => R:22 | X:1 => R:22 : R:22==R:22 unnecessary to merge!
merge X:1 & P:1 : X:1 => R:22 | P:1 => P:1 : merge into R:23
merge K:1 & L:1 : K:1 => R:23 | L:1 => R:23 : R:23==R:23 unnecessary to merge!
    
```

General Tree

- **K-ary tree**

- node & edge
- root
- 0 or K subtrees e.g. binary (2-ary) tree
- parent & children (0 or K)
- ancestor & descendant
- path & length
 - B-E-I, A-C-F : two paths of length 2
 - A-B-E-I-J: a path of length 4
- depth (cardinal) & level (ordinal)
- height (largest depth+1)
- leaf node & internal node





THANK YOU



上海交通大學

SHANGHAI JIAO TONG UNIVERSITY

Indexing



REVIEW

- **Primary (main) memory vs. secondary (peripheral) storage**
 - primary or main memory
 - e.g. **random access memory** (RAM), registers, cache, video memories
 - secondary or peripheral storage
 - e.g. **hard disk drives**, solid state drives, removable USB drives, CDs, DVDs
 - memory & storage access speed
 - RAM access speed is $10^5 \sim 10^6$ faster than disk drive access speed
- **Disk-based applications - minimize the number of disk accesses**
 - organize file structures
 - save information previously retrieved

Indexing



REVIEW

• Disk drive access

- disk drive architecture
- disk drive access
 - seek
 - preparatory rotation
 - data transfer

— disk drive access time (e.g.)

- disk drive (20G) = 10 platters ; platter = 16384 tracks
- track (128K) = 32 clusters = 256 sectors ; sector = 0.5K ; cluster = 4K
- track-to-track seek = 2.0ms ; random seek = 8.0ms ; rotation = 8.33ms (7200 RPM)
- **how long will it take to read a file of 1M (1024K = 8 tracks = 256 clusters)?**
- if the file fills 8 adjacent tracks : 122ms
- if the file fills 256 clusters spread randomly across the disk
 - $256 \times [8.0 + 8.33 \times (0.5 + 8/256)] = 256 \times (8.0 + 4.43) = 3182\text{ms}$



Indexing



REVIEW

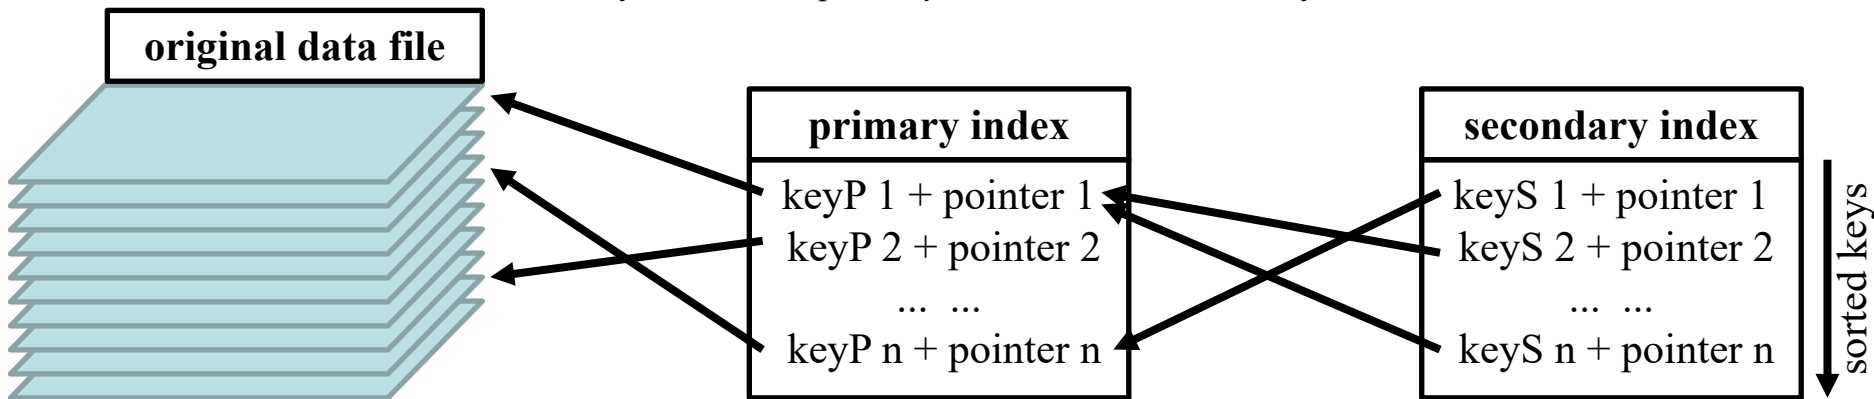
- **Disk drive access**
 - disk drive access time (e.g.)
 - disk drive (20G) = 10 platters ; platter = 16384 tracks
 - track (128K) = 32 clusters = 256 sectors ; sector = 0.5K ; cluster = 4K
 - track-to-track seek = 2.0ms ; random seek = 8.0ms ; rotation = 8.33ms (7200 RPM)
 - access a track : $8.0 + 8.33 \times (0.5 + 1) = 20.50\text{ms}$
 - access a sector : $8.0 + 8.33 \times (0.5 + 1/256) = 12.20\text{ms}$ (saves almost a half of track-access time)
 - access a byte : $8.0 + 8.33 \times (0.5 + 1/128K) = 12.17\text{ms}$ (save almost none of sector-access time)
 - **automatic reading/writing of an entire sector**
 - *buffering* a.k.a. *caching*
 - take/send additional information from/to disk to satisfy (potential) future requests
 - most operating systems maintain at least two buffers, one for input, one for output

Indexing



REVIEW


- **Index file - disk-based applications**
 - key sort & pointer methodology
 - **primary (key) index**
 - each record normally has a *unique identifier* called the primary key
 - relate each primary key value with a *pointer* to the actual record on disk
 - **secondary (key) index**
 - normally refer to the *primary index* instead of directly to the actual record on disk



Indexing



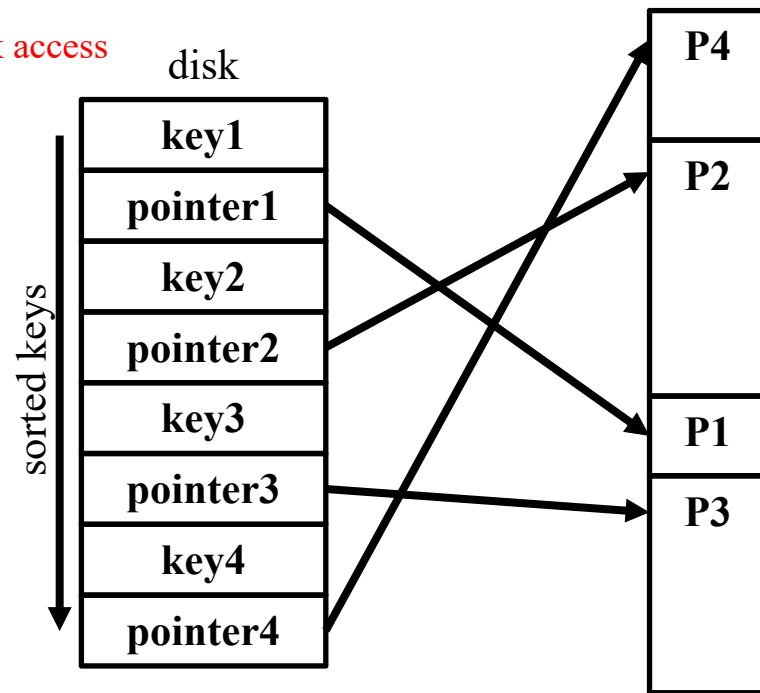
- **Linear indexing**
 - fixed-length key & pointer pairs



Indexing




- **Linear indexing**
 - fixed-length key & pointer pairs
 - searching involves a lot of inefficient disk access
 - automatic I/O of each entire block



Indexing

- **Linear indexing**

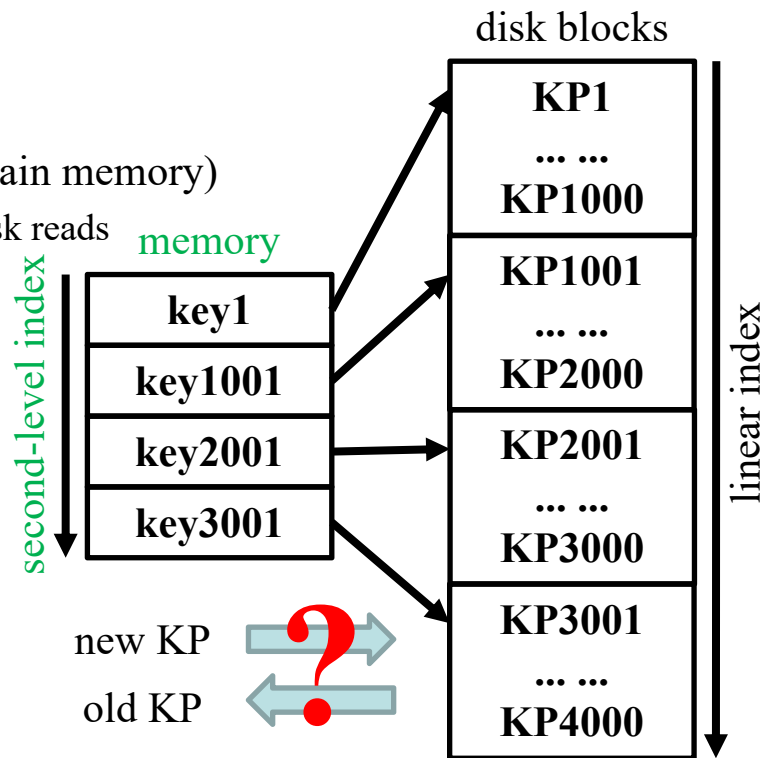
- fixed-length key & pointer pairs
 - automatic I/O of each entire block
- linear index & **second-level index** (in main memory)
 - accessing a record requires (only) two disk reads
 - one from the linear index file
 - one from the original data file



Indexing

- **Linear indexing**

- fixed-length key & pointer pairs
 - automatic I/O of each entire block
- linear index & **second-level index** (in main memory)
 - accessing a record requires (only) two disk reads
 - one from the linear index file
 - one from the original data file
- **updates to a linear index are expensive**
 - both insertion & removal are expensive
 - entire array contents might be shifted




Indexing



- **Linear indexing**

- linear index & **second-level index** (in main memory)
- *why updates to a linear index are expensive?*
 - perspective of *general tree*
 - **the tree is so rigid, denying flexible & efficient updates**



Indexing

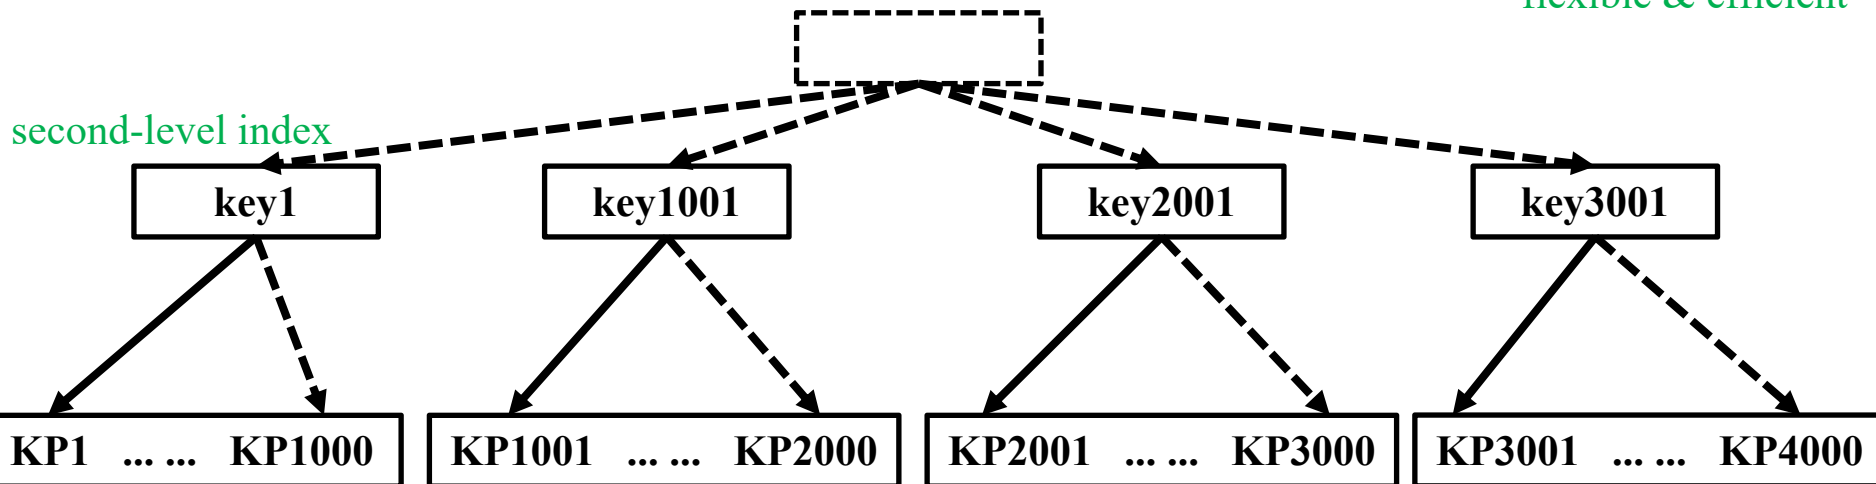


- **Linear indexing**

- linear index & **second-level index** (in main memory)
- *why updates to a linear index are expensive?*
 - perspective of *general tree*
 - *the tree is so rigid, denying flexible & efficient updates*



B-Tree
flexible & efficient



Indexing




- **B-Tree (of order m)**
 - root is either a leaf or has at least two children
 - internal node (except for root) has between $\lceil m/2 \rceil$ & m children
 - every node in the tree will be full at least to a certain minimum percentage
 - always *height balanced*, namely all leaves are at the same level
 - resonable distribution of records
 - internal node has *multiple keys* that separate its children
 - update & search operations affect only a few disk blocks & hence involve less disk I/O
 - keep related records i.e. records with similar key values on the same disk block, which helps to minimize I/O due to *locality of reference*
 - B-tree node is normally equivalent to a *disk block*

Indexing



- **B-Tree (of order m)**
 - internal node (except for root) has between $\lceil m/2 \rceil$ & m children
 - always *height balanced*, namely all leaves are at the same level
 - internal node has *multiple keys* that separate its children

B-tree of order four




Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - internal node (except for root) has between $\lceil m/2 \rceil$ & m children
 - always *height balanced*, namely all leaves are at the same level
 - internal node has *multiple keys* that separate its children
 - store records (or record pointers, primary key pointers) only at leaves

B⁺-tree of order four



leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - *binary search* may be applied within each node (block)
 - e.g. search 59

B⁺-tree of order four



search 59 : $31 \leq 59$ to second



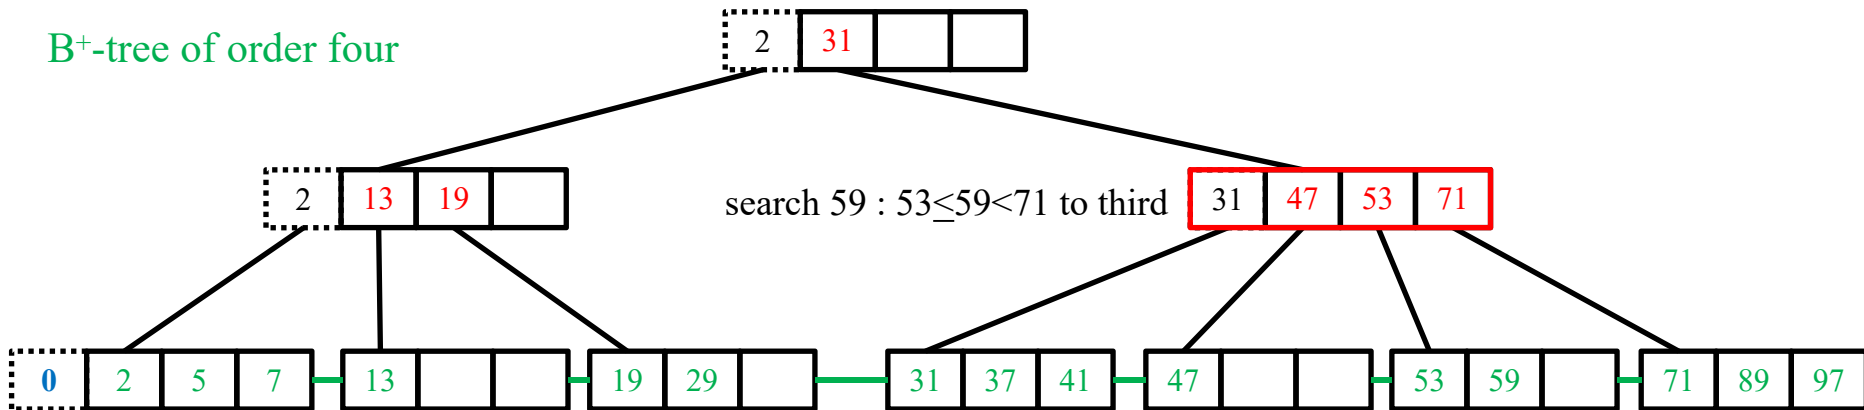
leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - *binary search* may be applied within each node (block)
 - e.g. search 59

B⁺-tree of order four




leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - *binary search* may be applied within each node (block)
 - e.g. search 59

B⁺-tree of order four




leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - *binary search* may be applied within each node (block)
 - e.g. search 59

B⁺-tree of order four



leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - *binary search* may be applied within each node (block)
 - e.g. search 59, search 17

B⁺-tree of order four



search 17 : $2 \leq 17 < 31$ to first



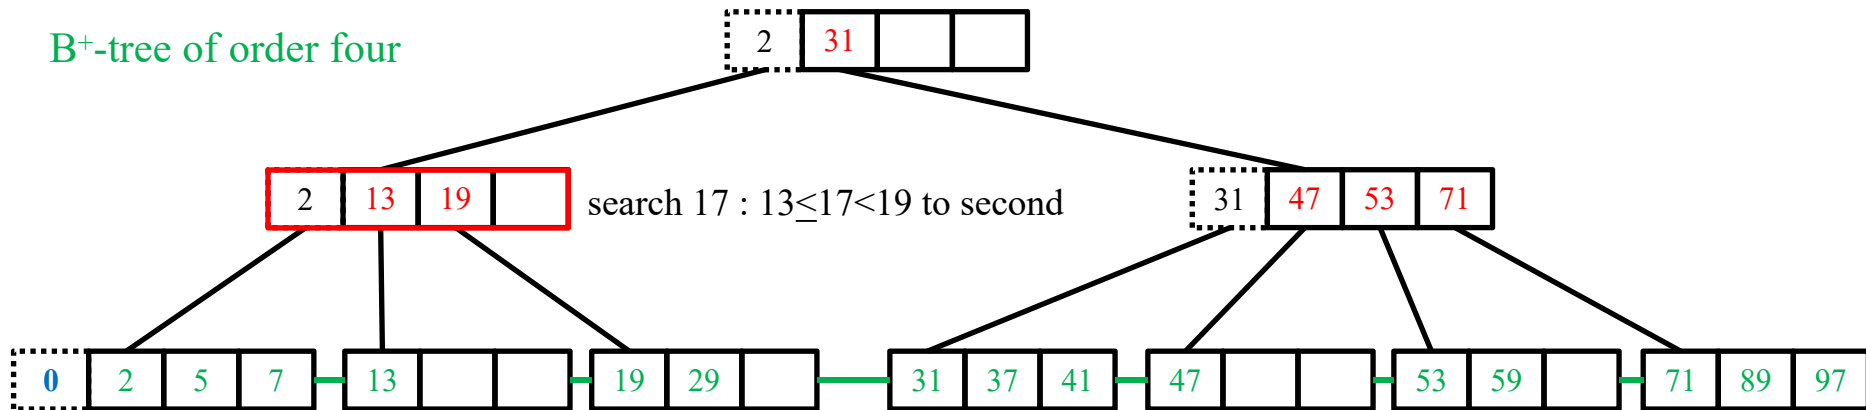
leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - *binary search* may be applied within each node (block)
 - e.g. search 59, search 17

B⁺-tree of order four




leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - *binary search* may be applied within each node (block)
 - e.g. search 59, search 17

B⁺-tree of order four




leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - *binary search* may be applied within each node (block)
 - e.g. search 59, search 17

B⁺-tree of order four



leaves are normally linked together (once first found, others probed sequentially for range query)

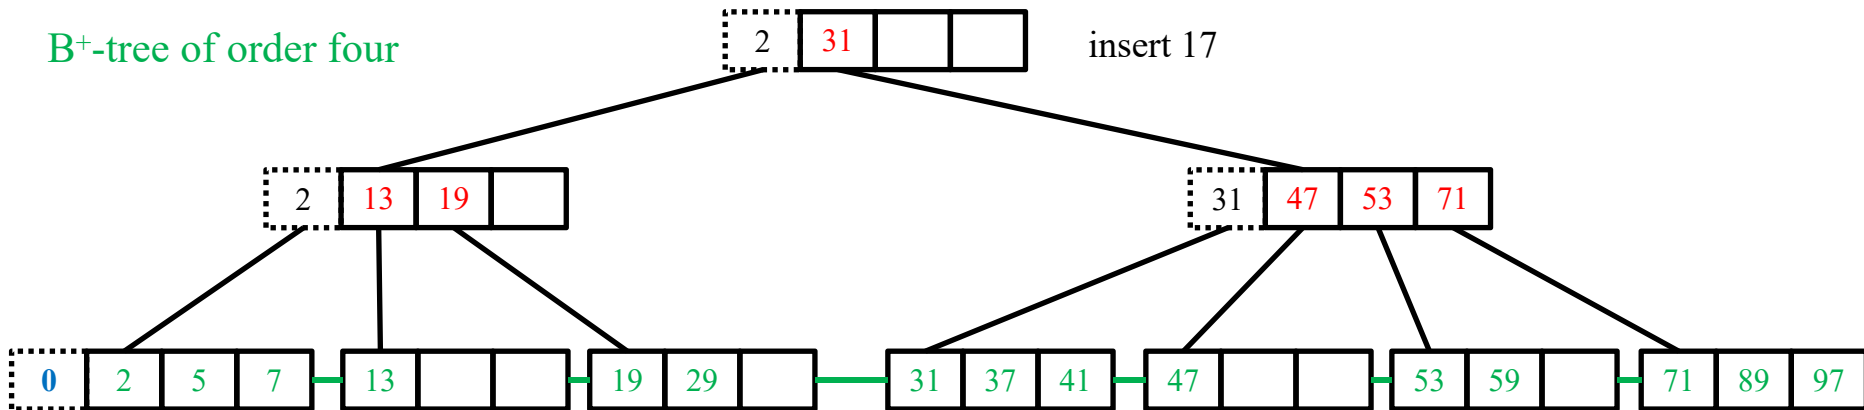
Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected

B⁺-tree of order four

insert 17




leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected

B⁺-tree of order four



leaves are normally linked together (once first found, others probed sequentially for range query)

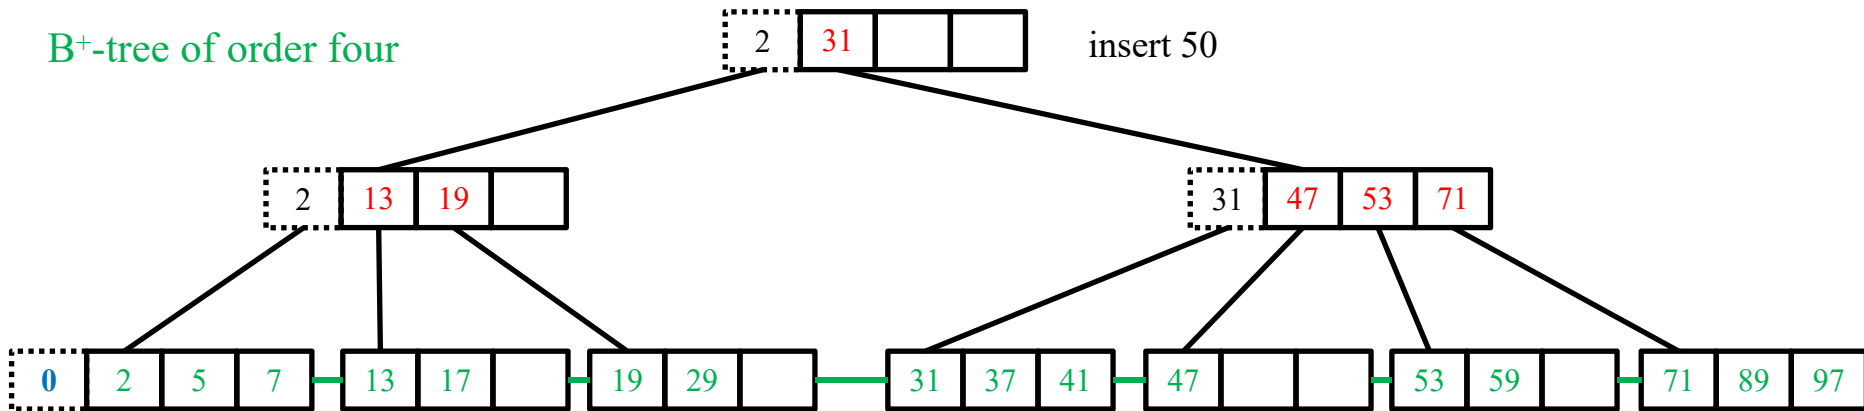
Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected

B⁺-tree of order four

insert 50




leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected

B⁺-tree of order four

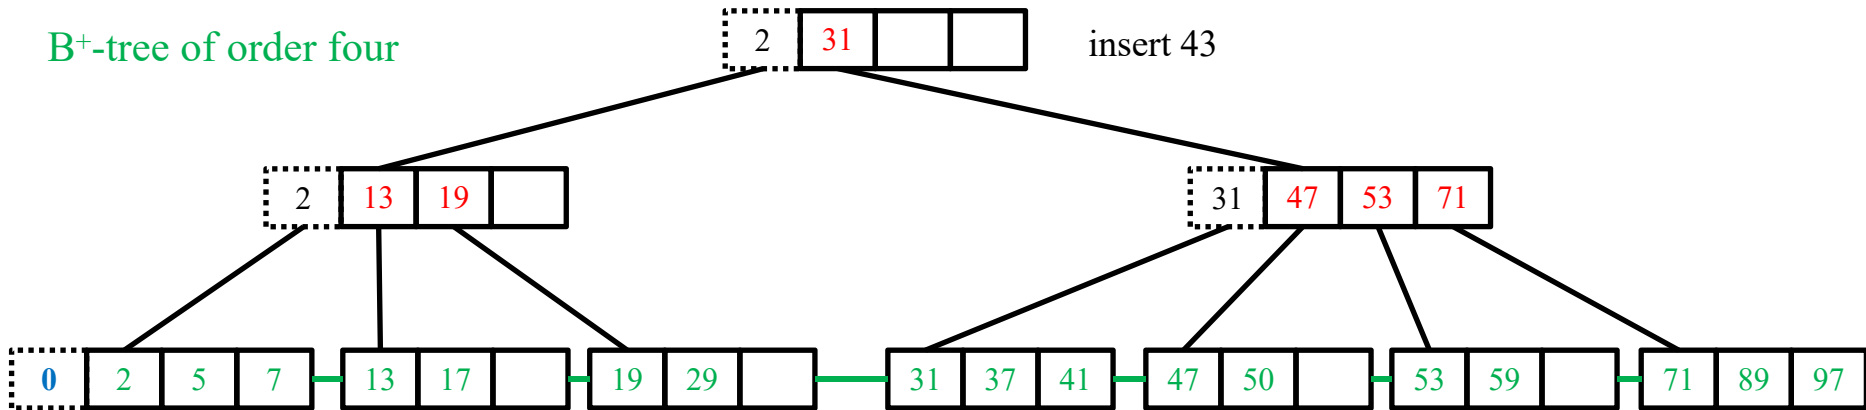


leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected
 - if leaf is full, perform **splitting-promotion**
 - split the node by *middle* & promote *middle* to parent

B⁺-tree of order four




leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected
 - if leaf is full, perform **splitting-promotion**
 - split the node by *middle* & promote *middle* to parent

B⁺-tree of order four




leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected
 - if leaf is full, perform **splitting-promotion**
 - split the node by *middle* & promote *middle* to parent

B⁺-tree of order four

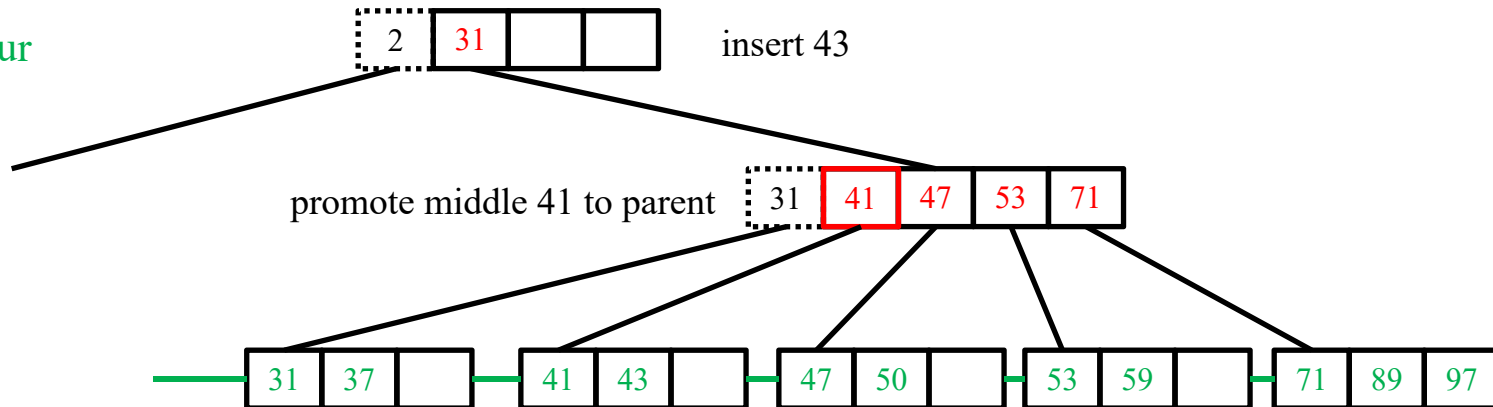


Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected
 - if leaf is full, perform **splitting-promotion**
 - split the node by *middle* & promote *middle* to parent

B⁺-tree of order four




leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected
 - if leaf is full, perform **splitting-promotion**
 - split the node by *middle* & promote *middle* to parent

B⁺-tree of order four



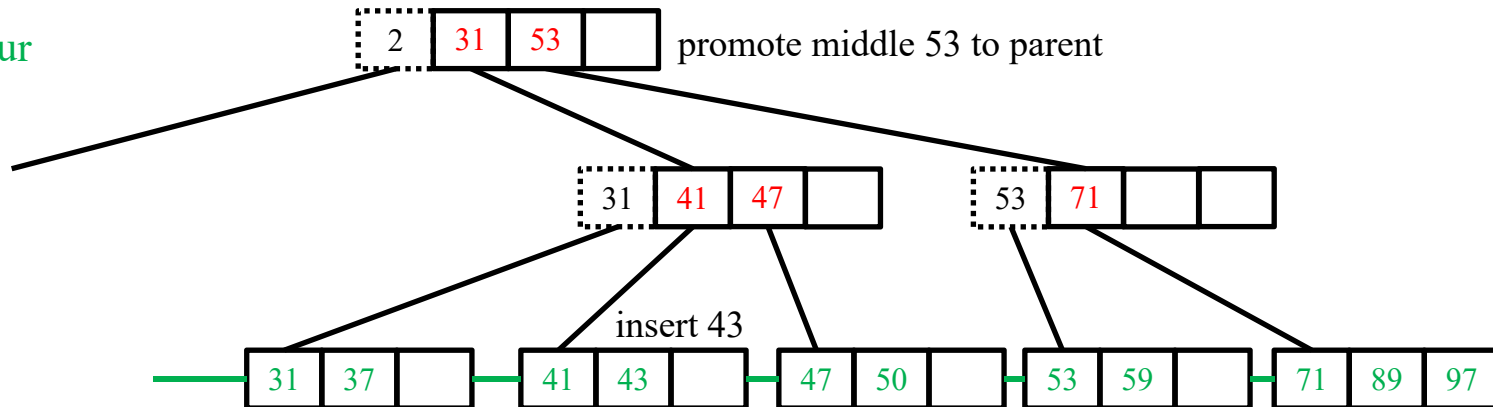
leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected
 - if leaf is full, perform **splitting-promotion**
 - split the node by *middle* & promote *middle* to parent

B⁺-tree of order four



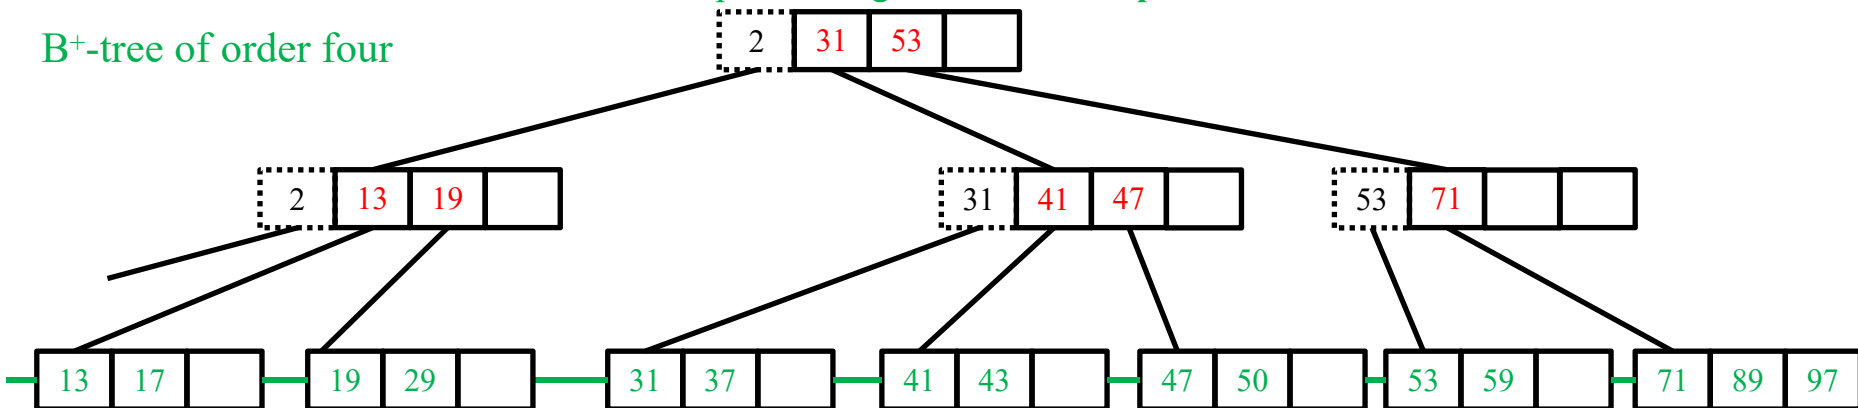
leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert
 - if leaf is not full, just insert into the leaf; no other B⁺-tree node is affected
 - if leaf is full, perform **splitting-promotion**
 - such insertion process is guaranteed to **keep all nodes at least half full**

B⁺-tree of order four



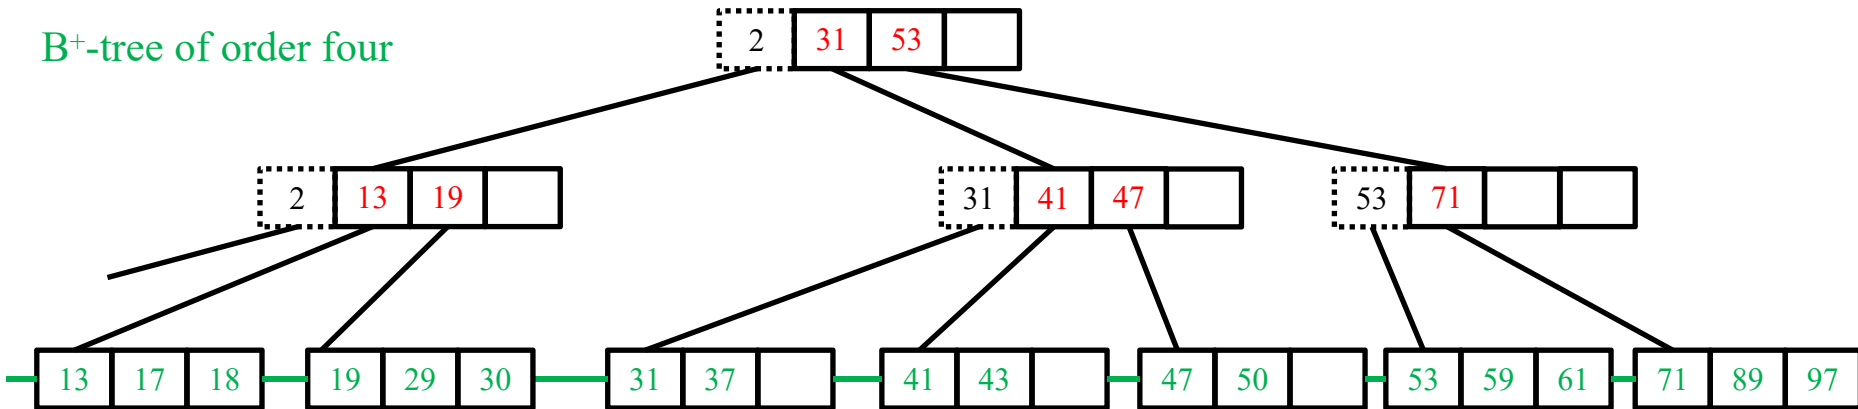
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - suppose all nodes are intended to be **kept at least half full** (similar logic applies for other minimum thresholds)

B⁺-tree of order four



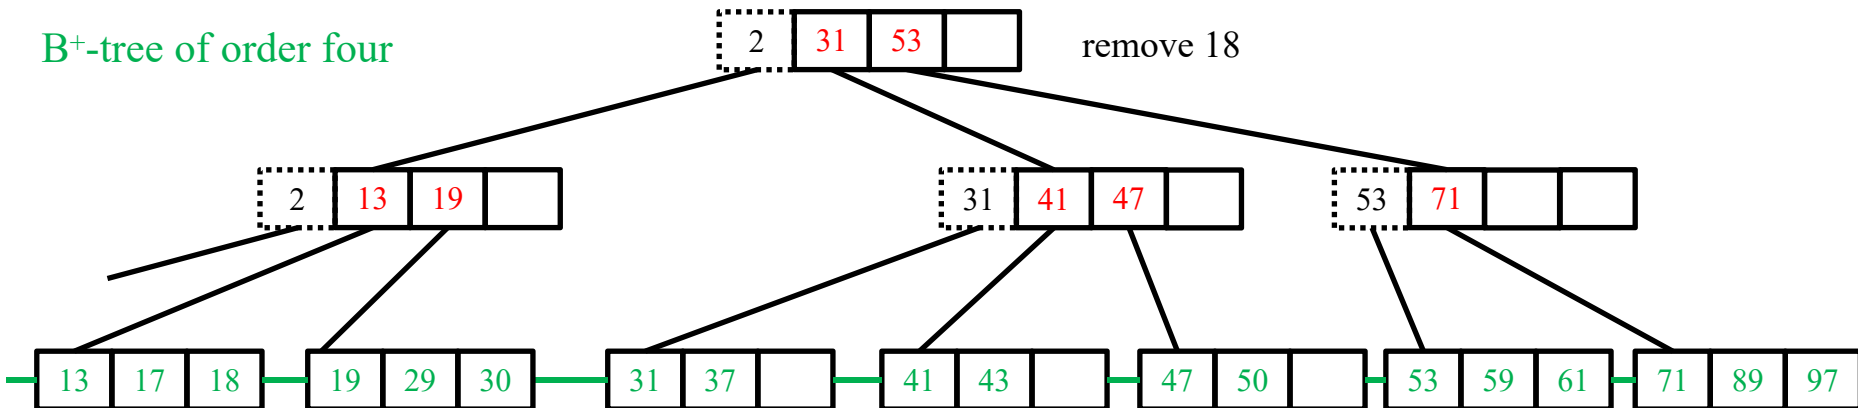
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf

B⁺-tree of order four



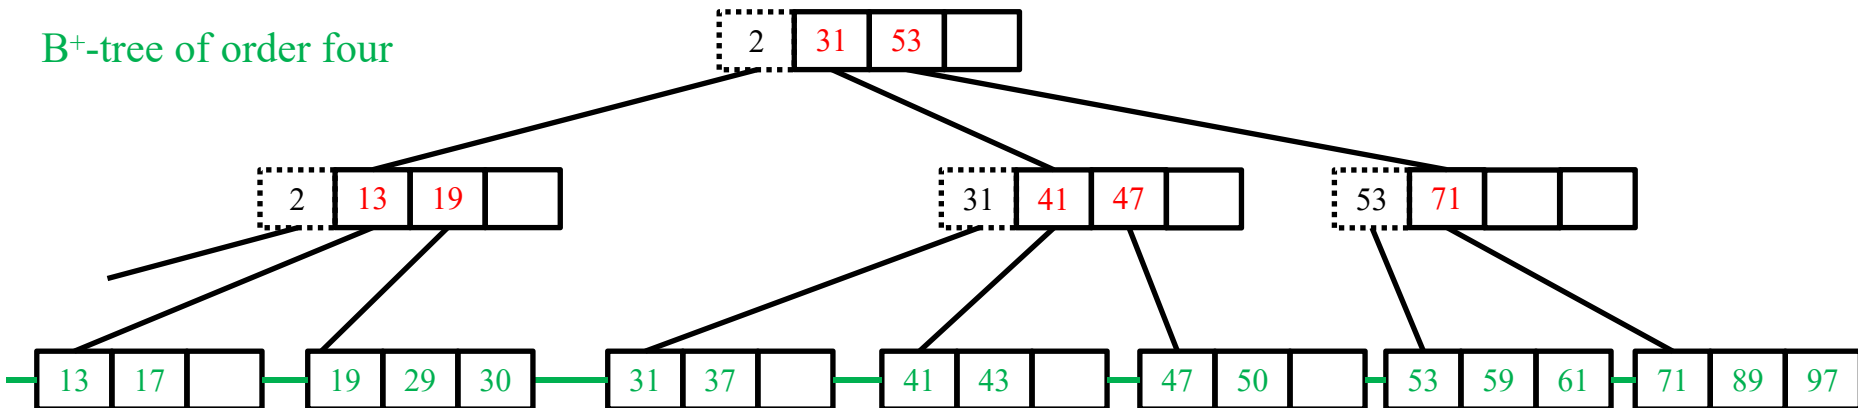
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf

B⁺-tree of order four



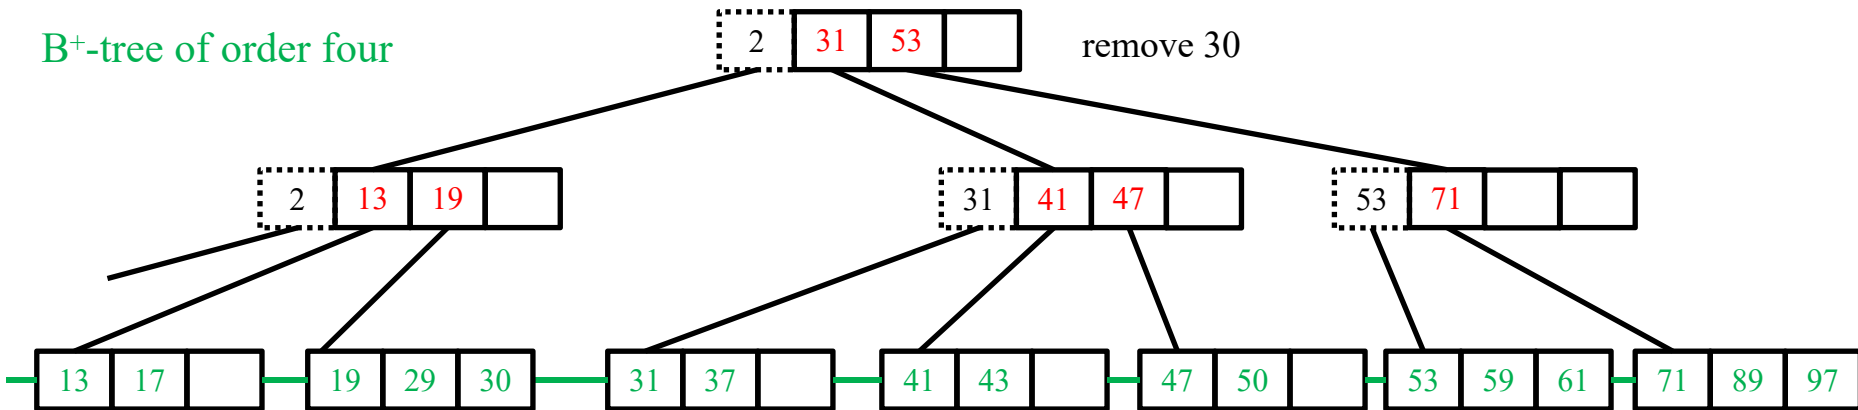
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf

B⁺-tree of order four



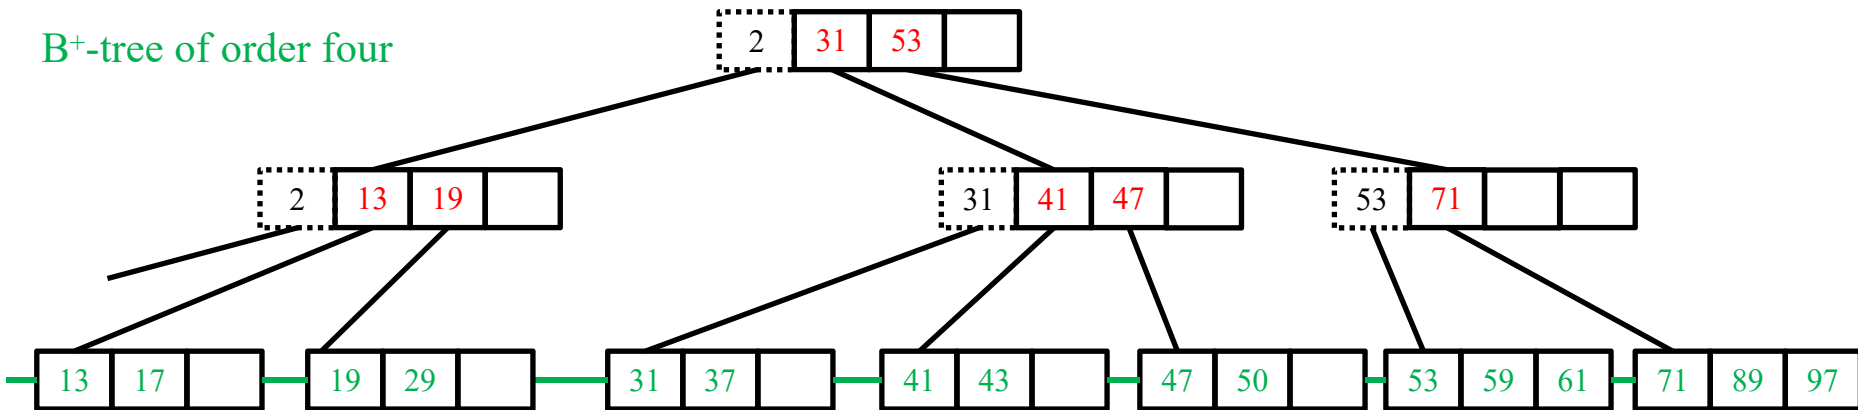
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf

B⁺-tree of order four



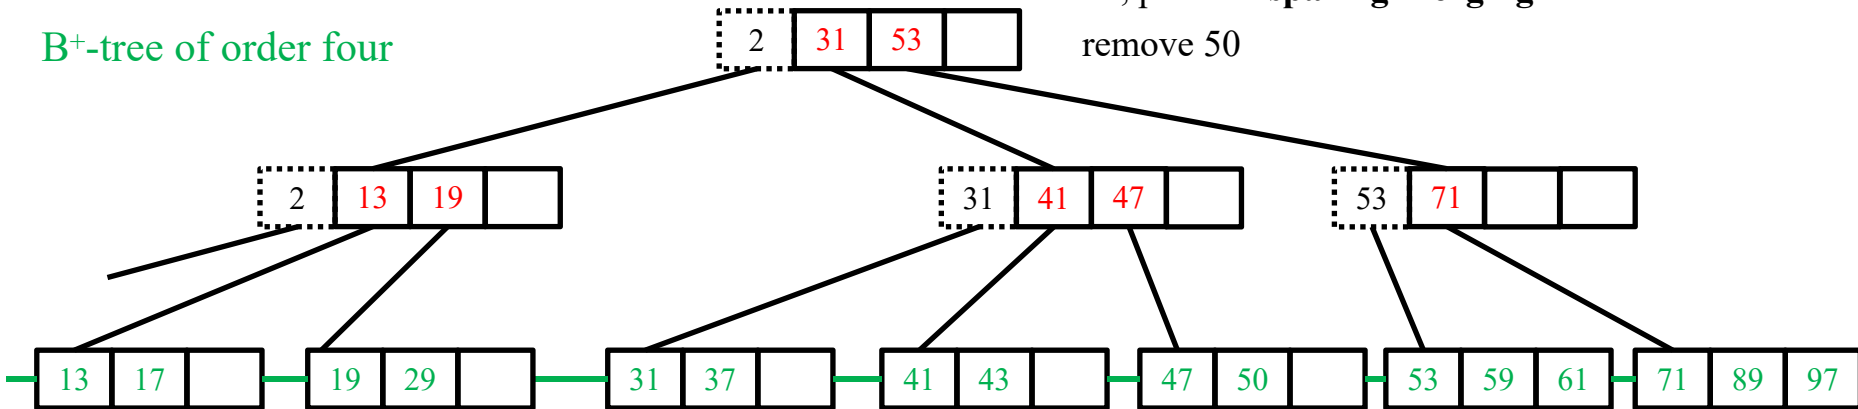
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four

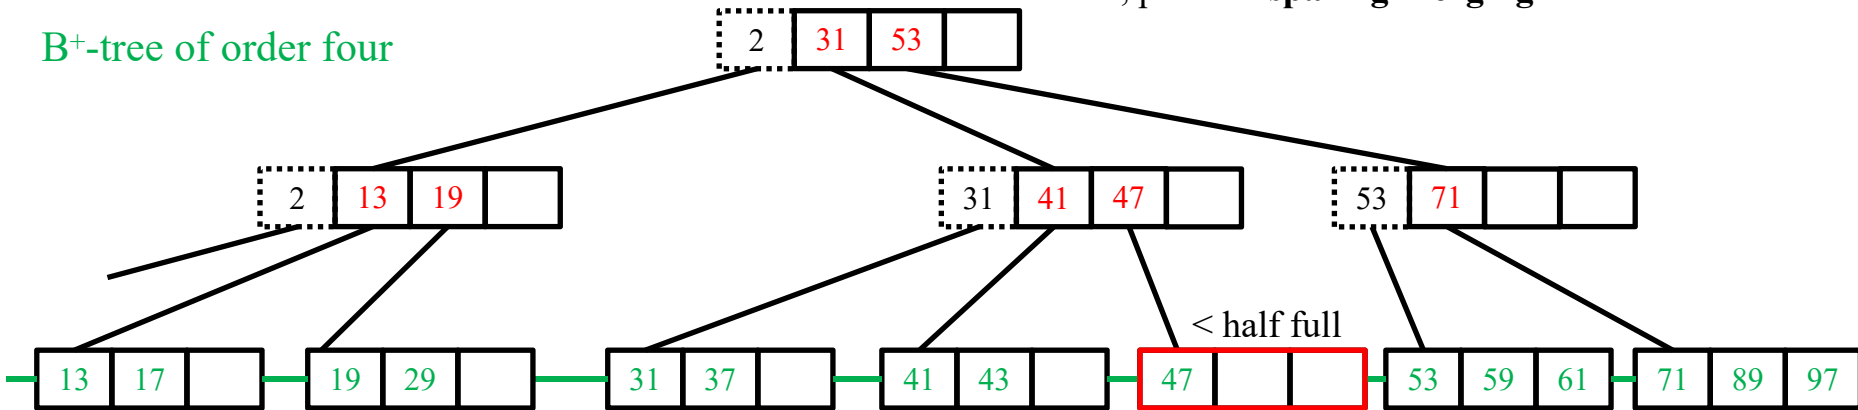


leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four

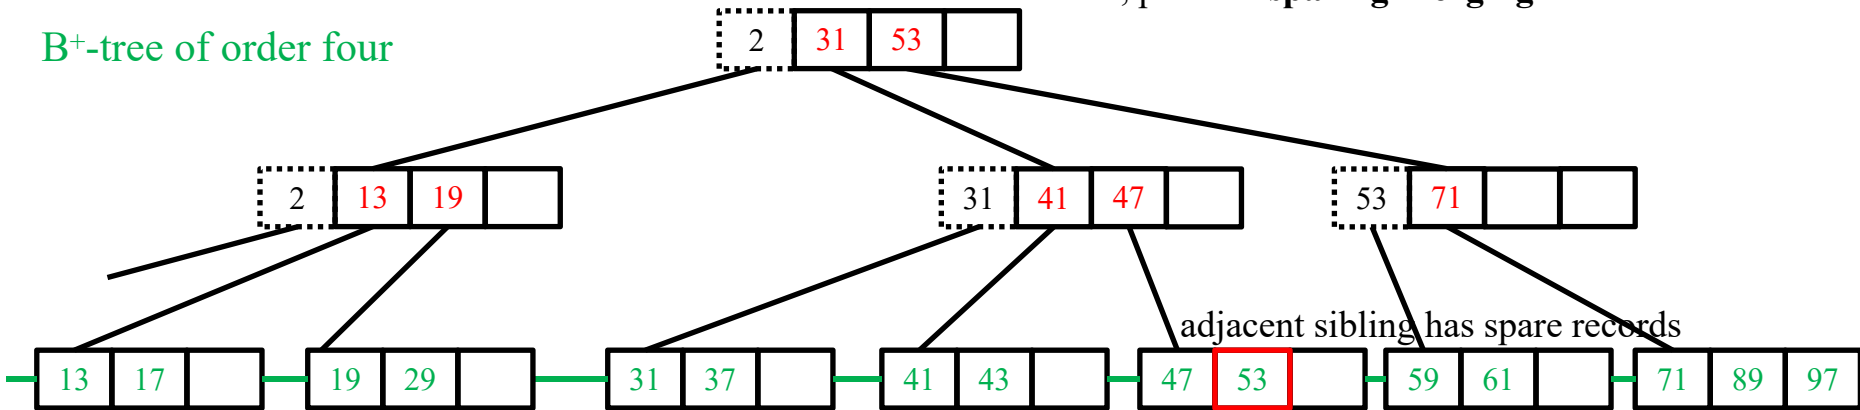


leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four



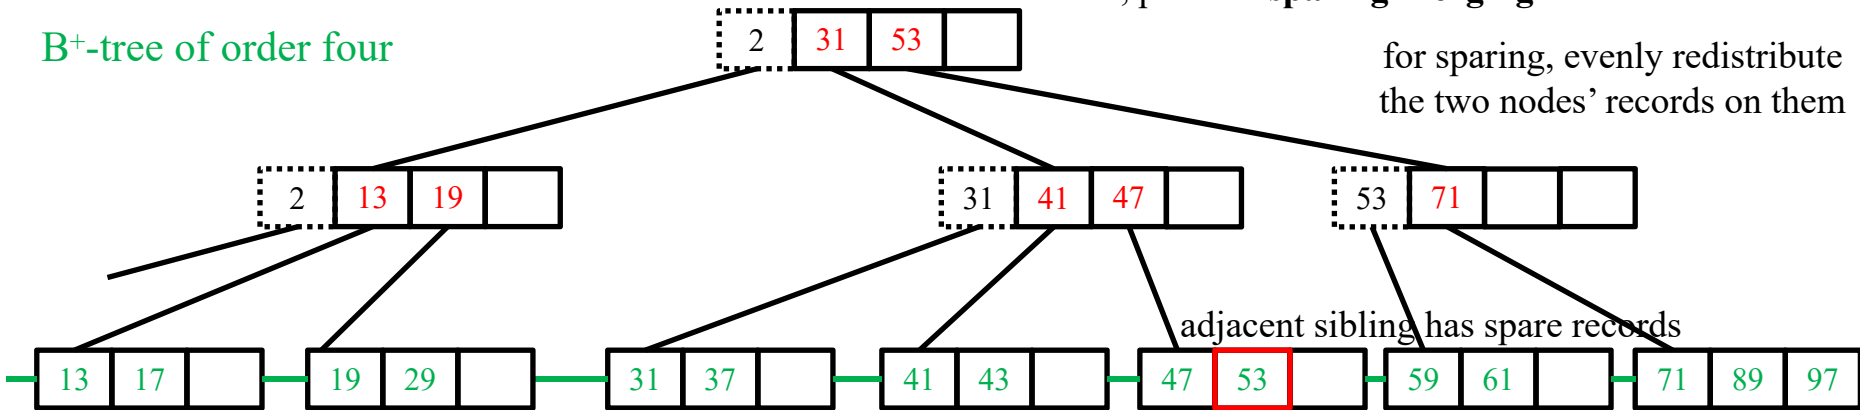
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**

- search
- insert - splitting-promotion
- remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four



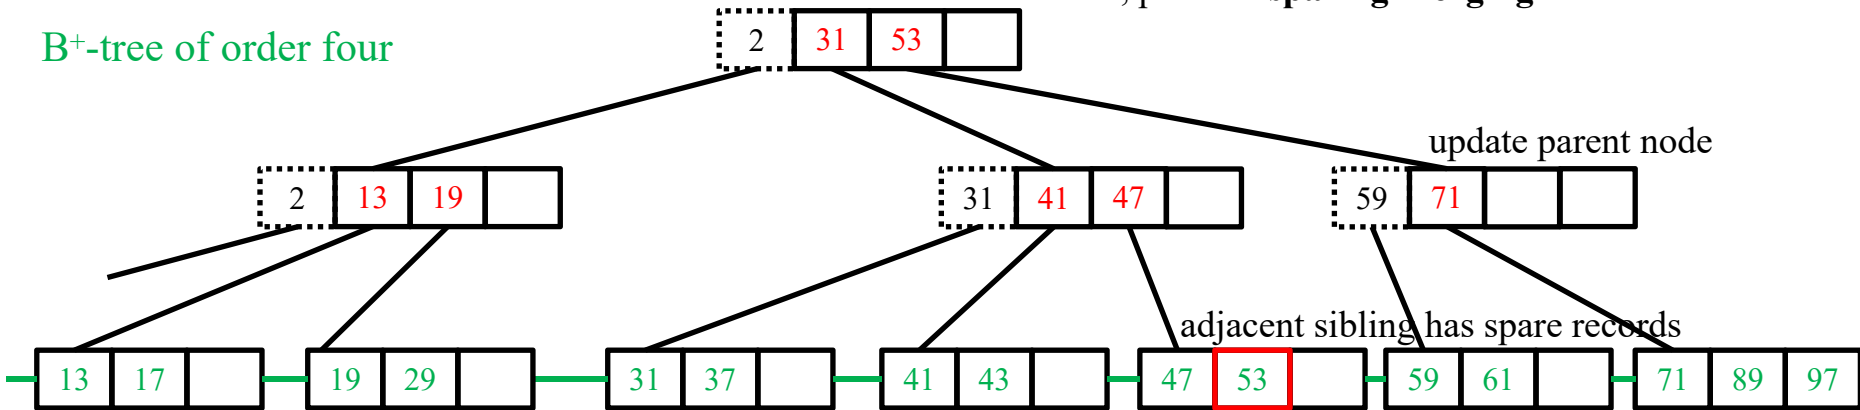
leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**

- search
- insert - splitting-promotion
- remove
 - if leaf can be at least half full after removal, just remove from the leaf
 - if leaf is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four



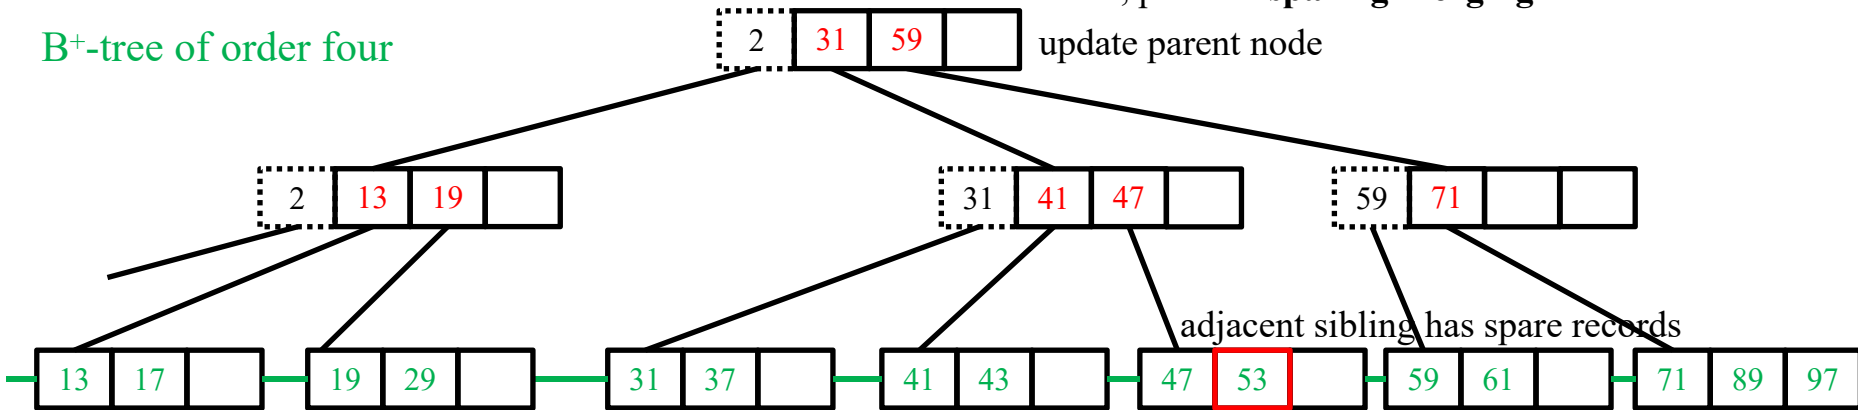
leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**

- search
- insert - splitting-promotion
- remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four



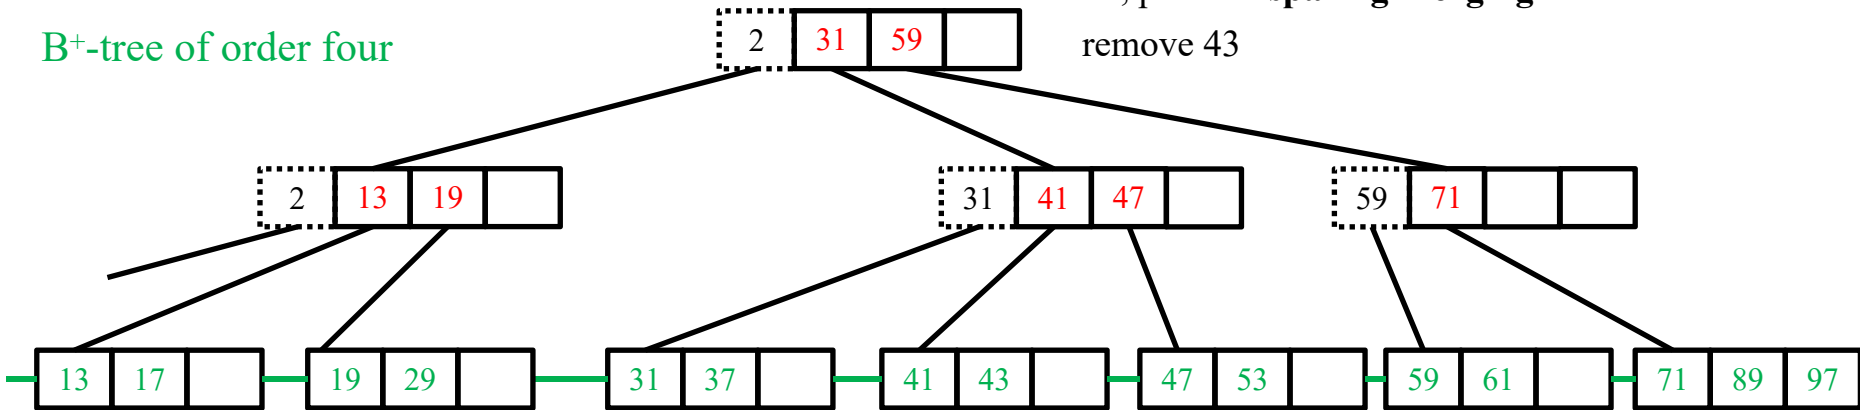
leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**

- search
- insert - splitting-promotion
- remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four

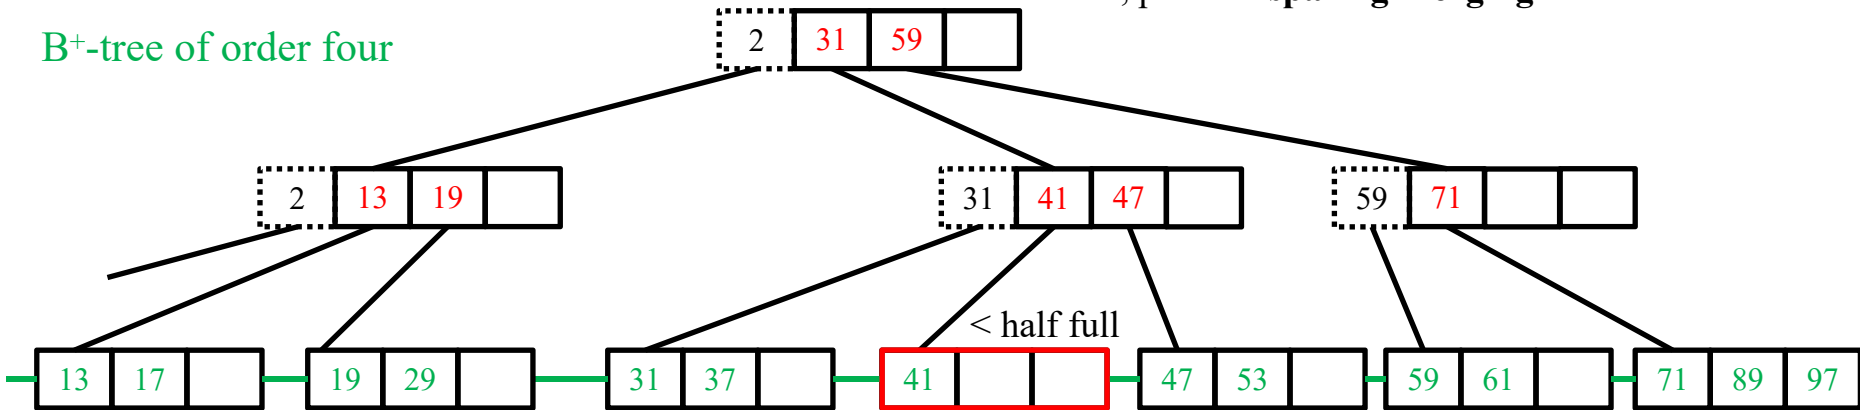


leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four

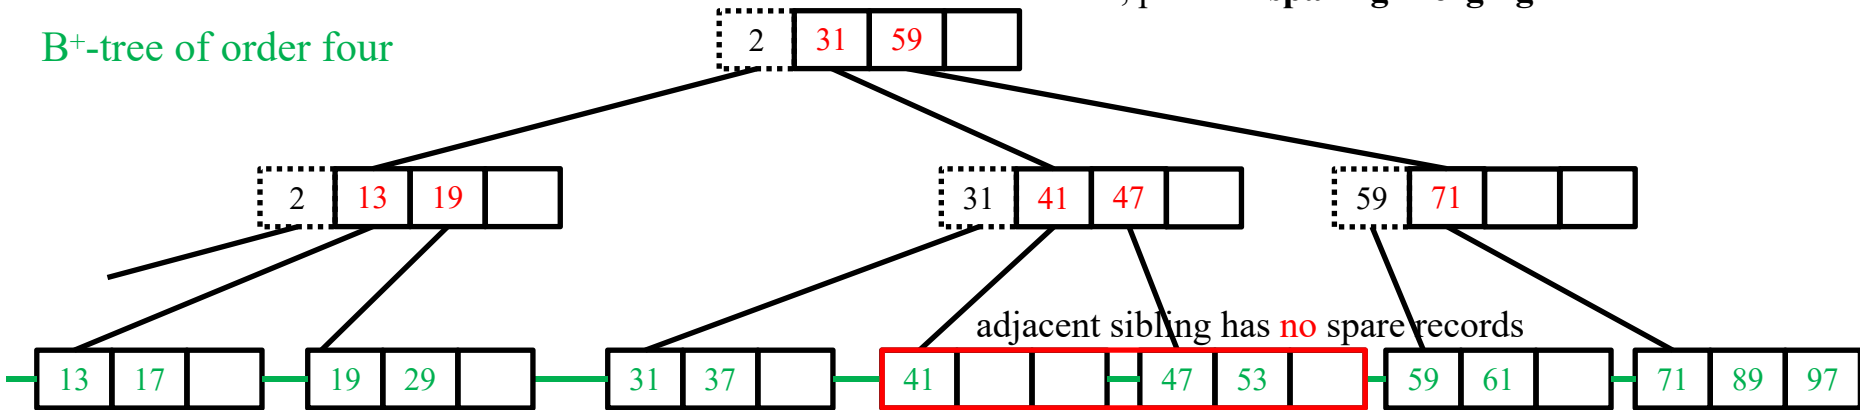


leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four

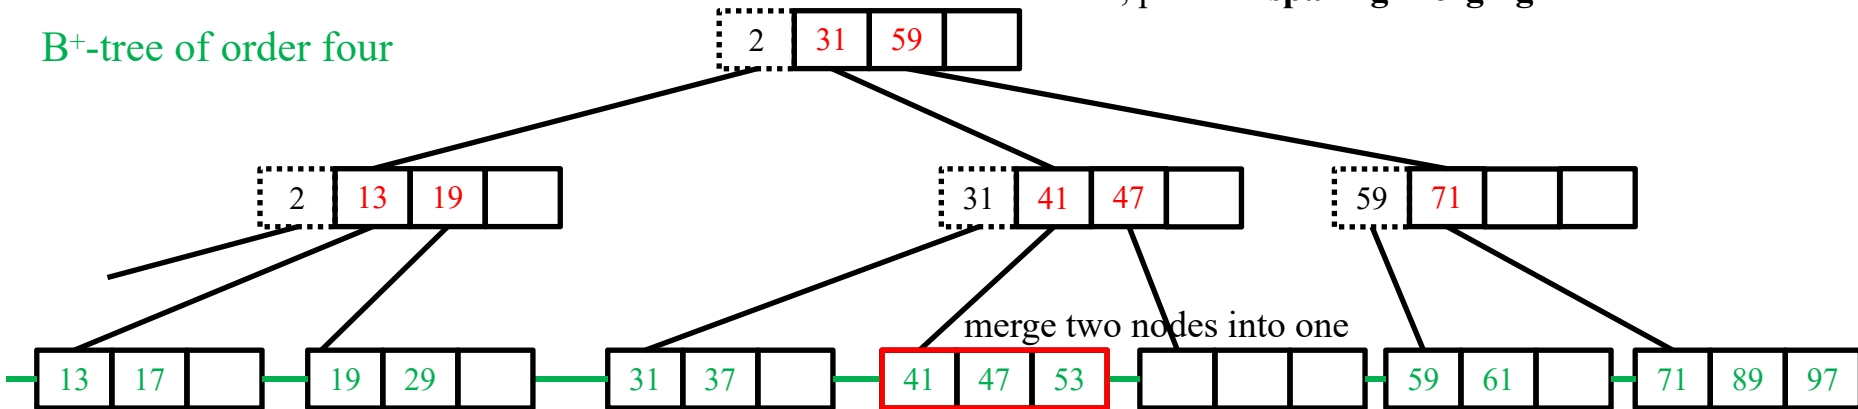


leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four

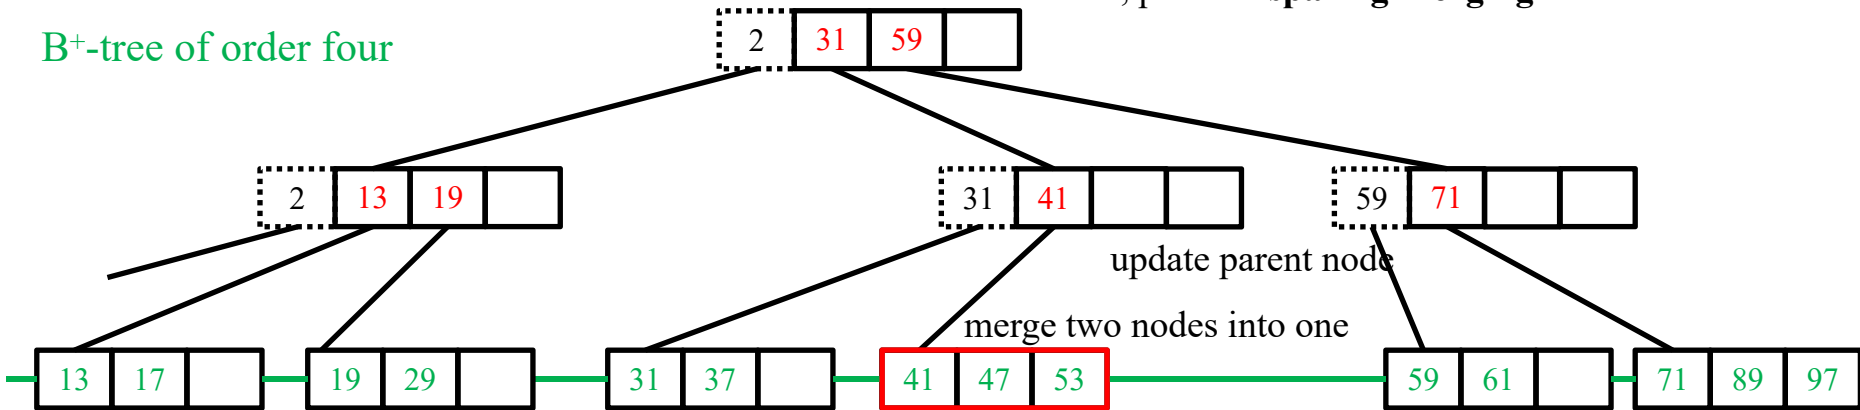


leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four



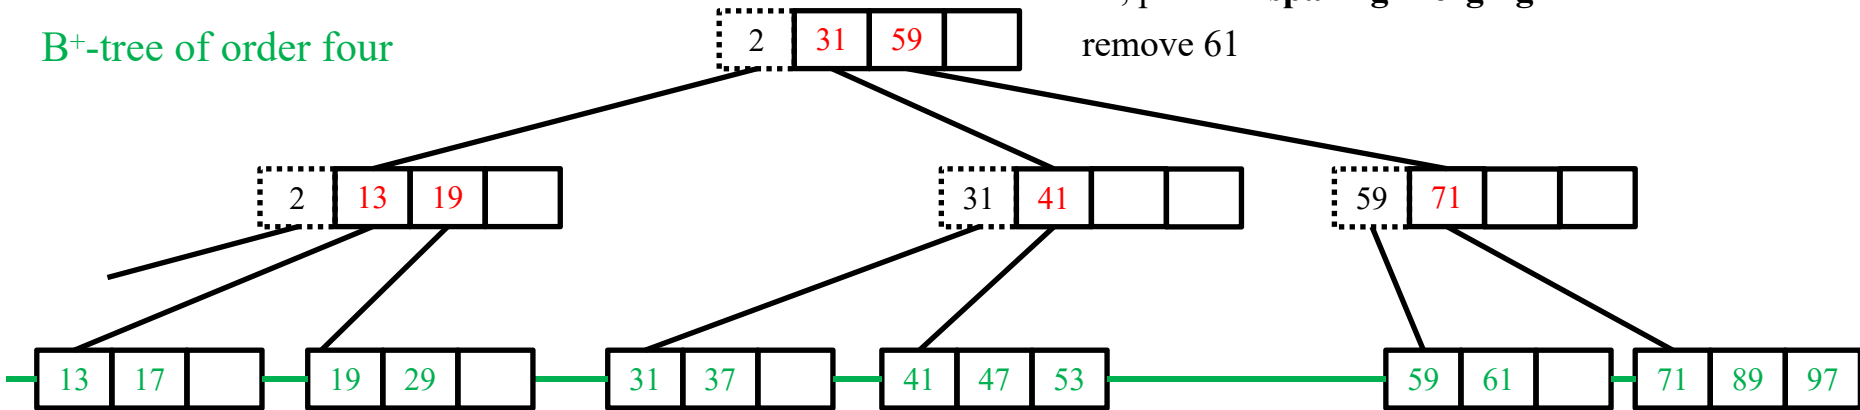
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**

- search
- insert - splitting-promotion
- remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four



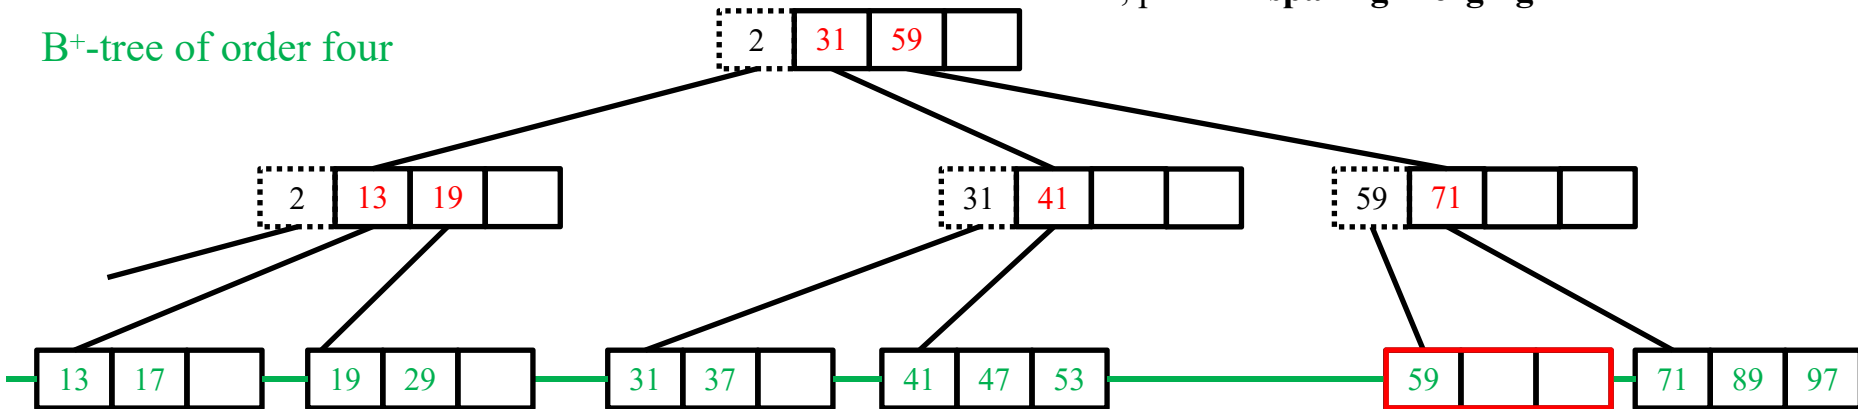
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four



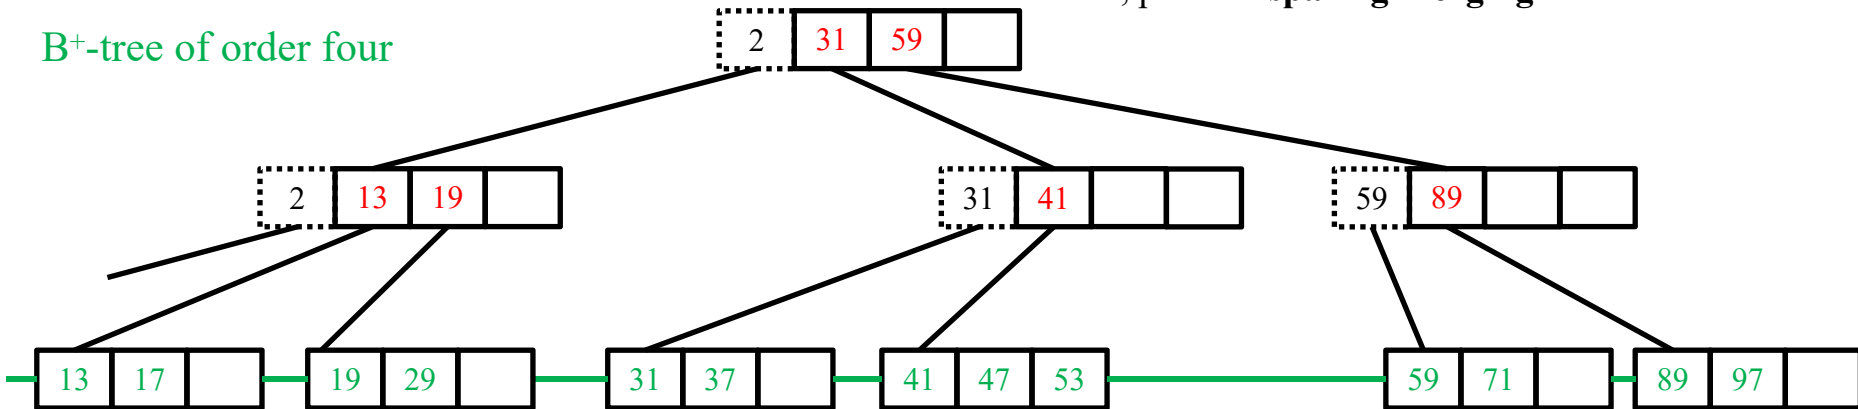
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four



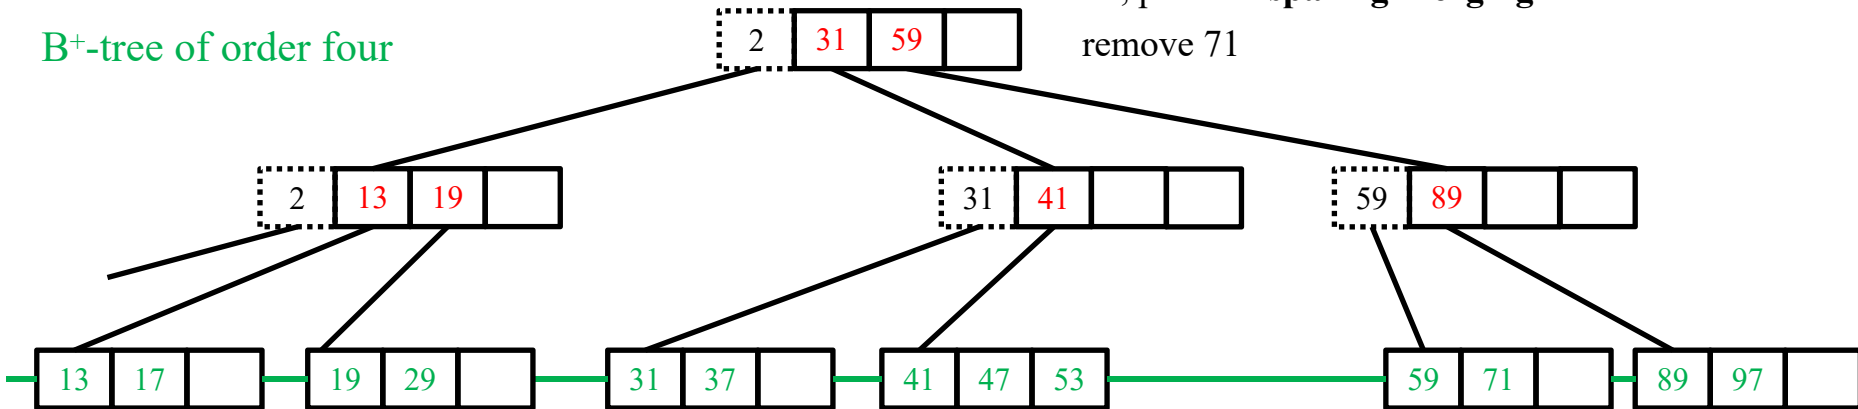
leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leaf is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four

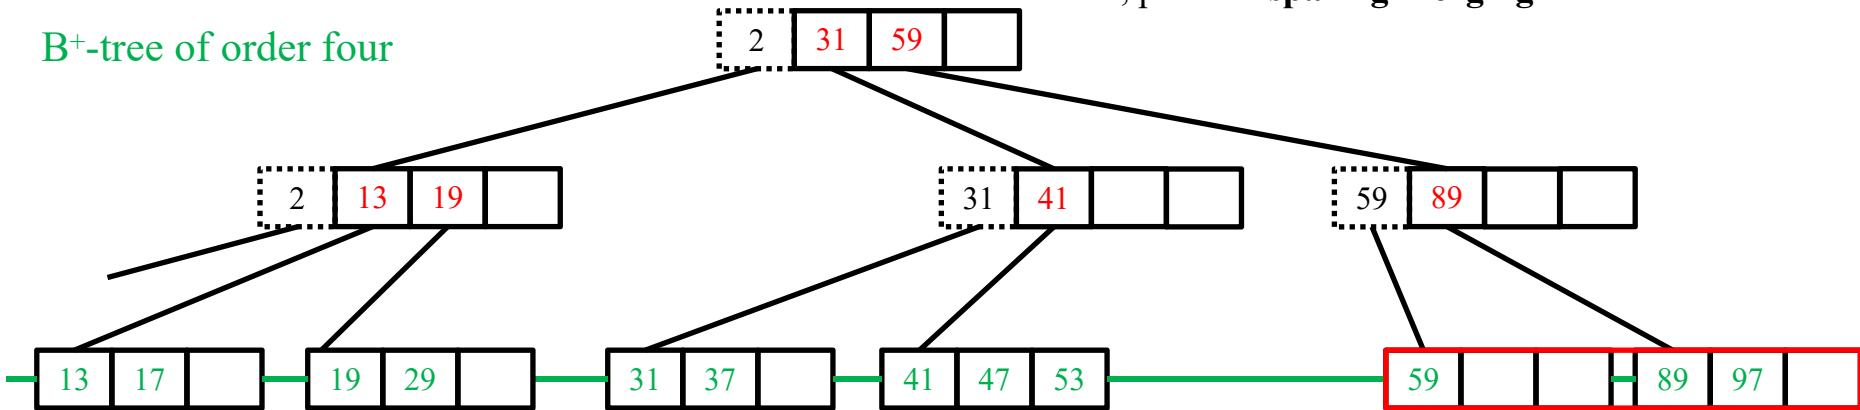


leaves are normally linked together (once first found, others probed sequentially for *range query*)

Indexing

- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four




leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leave
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four




leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B⁺-Tree (of order m) - actually adopted in practice**
 - search
 - insert - splitting-promotion
 - remove
 - if leave can be at least half full after removal, just remove from the leaf
 - if leave is less than half full after removal, perform **sparing-merging**

B⁺-tree of order four




leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing



- **B*-Tree (of order m) - actually adopted in practice**
 - insert - splitting-promotion
 - remove - sparing-merging
 - identical to B⁺-tree, except for rules of splitting & merging nodes
 - split two nodes into three instead of one into two
 - merge three nodes into two instead of two into one


B*-tree of order four



leaves are normally linked together (once first found, others probed sequentially for range query)

Indexing

- **B-Tree family (of order m) - B⁺/B^{*}-tree actually adopted in practice**
 - internal node (except for root) has between $\lceil m/2 \rceil$ & m children
 - always *height balanced*, namely all leaves are at the same level
 - internal node has *multiple keys* that separate its children
 - store records (or record pointers, primary key pointers) only at leaves (B⁺/B^{*}-tree)
 - insertion process is guaranteed to **keep all nodes at least half full**
 - order m is large (e.g. 100) in practice, so overhead for internal nodes is very low



leaves are normally linked together (once first found, others probed sequentially for range query)



THANK YOU



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY